

# Related Python Code

## 1 KNN

```
import numpy as np
import pdb

"""
This code was based off of code from cs231n at Stanford University,
and modified for ECE C147/C247 at UCLA.
"""

class KNN(object):

    def __init__(self):
        pass

    def train(self, X, y):
        """
        Inputs:
        - X is a numpy array of size (num_examples, D)
        - y is a numpy array of size (num_examples, )
        """
        self.X_train = X
        self.y_train = y

    def compute_distances(self, X, norm=None):
        """
        Compute the distance between each test point in X and each training point
        in self.X_train.

        Inputs:
        - X: A numpy array of shape (num_test, D) containing test data.
        - norm: the function with which the norm is taken.

        Returns:
        - dists: A numpy array of shape (num_test, num_train) where dists[i, j]
            is the Euclidean distance between the ith test point and the jth training
            point.
        """
        if norm is None:
            norm = lambda x: np.sqrt(np.sum(x**2))
            #norm = 2

        num_test = X.shape[0]
        num_train = self.X_train.shape[0]
```

```

dists = np.zeros((num_test, num_train))
for i in np.arange(num_test):
    for j in np.arange(num_train):
        # ===== #
        # YOUR CODE HERE:
        # Compute the distance between the ith test point and the jth
        # training point using norm(), and store the result in dists[i, j].
        # ===== #

        dists[i,j] = norm(X[i] - self.X_train[j])

        # ===== #
        # END YOUR CODE HERE
        # ===== #

return dists

def compute_L2_distances_vectorized(self, X):
    """
    Compute the distance between each test point in X and each training point
    in self.X_train WITHOUT using any for loops.

    Inputs:
    - X: A numpy array of shape (num_test, D) containing test data.

    Returns:
    - dists: A numpy array of shape (num_test, num_train) where dists[i, j]
      is the Euclidean distance between the ith test point and the jth training
      point.
    """
    num_test = X.shape[0]
    num_train = self.X_train.shape[0]
    dists = np.zeros((num_test, num_train))

    # ===== #
    # YOUR CODE HERE:
    # Compute the L2 distance between the ith test point and the jth
    # training point and store the result in dists[i, j]. You may
    # NOT use a for loop (or list comprehension). You may only use
    # numpy operations.
    #
    # HINT: use broadcasting. If you have a shape (N,1) array and
    # a shape (M,) array, adding them together produces a shape (N, M)
    # array.
    # ===== #

    xsq = np.square(X).sum(axis=1).reshape(num_test, 1)
    ysq = np.square(self.X_train).sum(axis=1)
    xy2 = 2*np.dot(X, (self.X_train.T))
    dists = np.sqrt(xsq + ysq - xy2)

    # ===== #
    # END YOUR CODE HERE

```

```

# ===== #

return dists

def predict_labels(self, dists, k=1):
    """
    Given a matrix of distances between test points and training points,
    predict a label for each test point.

    Inputs:
    - dists: A numpy array of shape (num_test, num_train) where dists[i, j]
      gives the distance between the ith test point and the jth training point.

    Returns:
    - y: A numpy array of shape (num_test,) containing predicted labels for the
      test data, where y[i] is the predicted label for the test point X[i].
    """
    num_test = dists.shape[0]
    y_pred = np.zeros(num_test)
    for i in np.arange(num_test):
        # A list of length k storing the labels of the k nearest neighbors to
        # the ith test point.
        closest_y = []
        # ===== #
        # YOUR CODE HERE:
        # Use the distances to calculate and then store the labels of
        # the k-nearest neighbors to the ith test point. The function
        # numpy.argsort may be useful.
        #
        # After doing this, find the most common label of the k-nearest
        # neighbors. Store the predicted label of the ith training example
        # as y_pred[i]. Break ties by choosing the smaller label.
        # ===== #
        idx = np.argsort(dists[i,:])
        near_x = idx[0:k]
        near_y = self.y_train[near_x]
        max_idx = np.argmax(np.bincount(near_y))
        y_pred[i] = np.amin(max_idx) # breaking ties by selecting the minimal index

        # ===== #
        # END YOUR CODE HERE
        # ===== #

    return y_pred

```

## 2 SVM

```
import numpy as np
import pdb

"""
This code was based off of code from cs231n at Stanford University,
and modified for ECE C147/C247 at UCLA.
"""

class SVM(object):

    def __init__(self, dims=[10, 3073]):
        self.init_weights(dims=dims)

    def init_weights(self, dims):
        """
        Initializes the weight matrix of the SVM. Note that it has shape (C, D)
        where C is the number of classes and D is the feature size.
        """
        self.W = np.random.normal(size=dims)

    def loss(self, X, y):
        """
        Calculates the SVM loss.

        Inputs have dimension D, there are C classes, and we operate on minibatches
        of N examples.

        Inputs:
        - X: A numpy array of shape (N, D) containing a minibatch of data.
        - y: A numpy array of shape (N,) containing training labels; y[i] = c means
            that X[i] has label c, where 0 <= c < C.

        Returns a tuple of:
        - loss as single float
        """

        # compute the loss and the gradient
        num_classes = self.W.shape[0]
        num_train = X.shape[0]
        loss = 0.0

        for i in np.arange(num_train):
            # ===== #
            # YOUR CODE HERE:
            # Calculate the normalized SVM loss, and store it as 'loss'.
            # (That is, calculate the sum of the losses of all the training
            # set margins, and then normalize the loss by the number of
            # training examples.)
            # ===== #
            for j in range(num_classes):
```

```

        z_j = 0
        if y[i] != j:
            z_j = 1 + self.W[j].T.dot(X[i]) - self.W[y[i]].T.dot(X[i])
            loss += max(0, z_j)
    loss /= num_train
    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return loss

def loss_and_grad(self, X, y):
    """
    Same as self.loss(X, y), except that it also returns the gradient.

    Output: grad -- a matrix of the same dimensions as W containing
             the gradient of the loss with respect to W.
    """

    # compute the loss and the gradient
    num_classes = self.W.shape[0]
    num_train = X.shape[0]
    loss = 0.0
    grad = np.zeros_like(self.W)

    for i in np.arange(num_train):
        # ===== #
        # YOUR CODE HERE:
        # Calculate the SVM loss and the gradient. Store the gradient in
        # the variable grad.
        # ===== #
        # get loss
        for j in range(num_classes):
            z_j = 0
            if y[i] != j:
                z_j = 1 + self.W[j].T.dot(X[i]) - self.W[y[i]].T.dot(X[i])
                loss += max(0, z_j)
            # get gradients
            grad[j] += X[i] * (z_j > 0)
            grad[y[i]] -= X[i] * (z_j > 0)

        # ===== #
        # END YOUR CODE HERE
        # ===== #

    loss /= num_train
    grad /= num_train

    return loss, grad

def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):
    """
    sample a few random elements and only return numerical

```

```

in these dimensions.
"""

for i in np.arange(num_checks):
    ix = tuple([np.random.randint(m) for m in self.W.shape])

    oldval = self.W[ix]
    self.W[ix] = oldval + h # increment by h
    fxph = self.loss(X, y)
    self.W[ix] = oldval - h # decrement by h
    fxmh = self.loss(X, y) # evaluate f(x - h)
    self.W[ix] = oldval # reset

    grad_numerical = (fxph - fxmh) / (2 * h)
    grad_analytic = your_grad[ix]
    rel_error = abs(grad_numerical - grad_analytic) / \
        (abs(grad_numerical) + abs(grad_analytic))
    print('numerical: %f analytic: %f, relative error: %e' %
        (grad_numerical, grad_analytic, rel_error))

def fast_loss_and_grad(self, X, y):
    """
    A vectorized implementation of loss_and_grad. It shares the same
    inputs and ouptuts as loss_and_grad.
    """
    loss = 0.0
    grad = np.zeros(self.W.shape) # initialize the gradient as zero

    # ===== #
    # YOUR CODE HERE:
    # Calculate the SVM loss WITHOUT any for loops.
    # ===== #
    num_train = X.shape[0]
    aj = np.dot(X, self.W.T)
    ay = aj[np.arange(num_train), y]
    ay = np.resize(ay, (num_train, 1))

    # after this, the loss of ay will become 1
    z = np.maximum(0, 1 + aj - ay)
    z[np.arange(num_train), y] = 0
    loss = np.sum(z)/num_train
    # ===== #
    # END YOUR CODE HERE
    # ===== #

    # ===== #
    # YOUR CODE HERE:
    # Calculate the SVM grad WITHOUT any for loops.
    # ===== #
    z[z > 0] = 1
    z[np.arange(num_train), y] = -np.sum(z, axis=1)
    grad = np.dot(z.T, X) / num_train
    # ===== #

```

```

# END YOUR CODE HERE
# ===== #

return loss, grad

def train(self, X, y, learning_rate=1e-3, num_iters=100,
          batch_size=200, verbose=False):
    """
    Train this linear classifier using stochastic gradient descent.

    Inputs:
    - X: A numpy array of shape (N, D) containing training data; there are N
        training samples each of dimension D.
    - y: A numpy array of shape (N,) containing training labels; y[i] = c
        means that X[i] has label 0 ≤ c < C for C classes.
    - learning_rate: (float) learning rate for optimization.
    - num_iters: (integer) number of steps to take when optimizing
    - batch_size: (integer) number of training examples to use at each step.
    - verbose: (boolean) If true, print progress during optimization.

    Outputs:
    A list containing the value of the loss function at each training iteration.
    """
    num_train, dim = X.shape
    # assume y takes values 0...K-1 where K is number of classes
    num_classes = np.max(y) + 1

    # initializes the weights of self.W
    self.init_weights(dims=[np.max(y) + 1, X.shape[1]])

    # Run stochastic gradient descent to optimize W
    loss_history = []

    for it in np.arange(num_iters):
        X_batch = None
        y_batch = None

        # ===== #
        # YOUR CODE HERE:
        # Sample batch_size elements from the training data for use in
        # gradient descent. After sampling,
        # - X_batch should have shape: (dim, batch_size)
        # - y_batch should have shape: (batch_size,)
        # The indices should be randomly generated to reduce correlations
        # in the dataset. Use np.random.choice. It's okay to sample with
        # replacement.
        # ===== #
        idx = np.random.choice(len(X), size=batch_size, replace=False)
        X_batch = X[idx]
        y_batch = y[idx]
        # print(X_batch.shape, y_batch.shape)
        # ===== #
        # END YOUR CODE HERE

```

```

# ===== #

# evaluate loss and gradient
loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
loss_history.append(loss)

# ===== #
# YOUR CODE HERE:
#   Update the parameters, self.W, with a gradient step
# ===== #
self.W -= learning_rate * grad
# ===== #
# END YOUR CODE HERE
# ===== #

if verbose and it % 100 == 0:
    print('iteration {} / {}: loss {}'.format(it, num_iters, loss))

return loss_history

def predict(self, X):
    """
    Inputs:
    - X: N x D array of training data. Each row is a D-dimensional point.

    Returns:
    - y_pred: Predicted labels for the data in X. y_pred is a 1-dimensional
      array of length N, and each element is an integer giving the predicted
      class.
    """
    y_pred = np.zeros(X.shape[1])

    # ===== #
    # YOUR CODE HERE:
    #   Predict the labels given the training data with the parameter self.W.
    # ===== #
    y_pred = np.argmax(np.dot(self.W, X.T), axis=0)
    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return y_pred

```



### 3 Softmax

```
import numpy as np

class Softmax(object):

    def __init__(self, dims=[10, 3073]):
        self.init_weights(dims=dims)

    def init_weights(self, dims):
        """
        Initializes the weight matrix of the Softmax classifier.
        Note that it has shape (C, D) where C is the number of
        classes and D is the feature size.
        """
        self.W = np.random.normal(size=dims) * 0.0001

    def loss(self, X, y):
        """
        Calculates the softmax loss.

        Inputs have dimension D, there are C classes, and we operate on minibatches
        of N examples.

        Inputs:
        - X: A numpy array of shape (N, D) containing a minibatch of data.
        - y: A numpy array of shape (N,) containing training labels; y[i] = c means
            that X[i] has label c, where 0 ≤ c < C.

        Returns a tuple of:
        - loss as single float
        """

        # Initialize the loss to zero.
        loss = 0.0

        # ===== #
        # YOUR CODE HERE:
        # Calculate the normalized softmax loss. Store it as the variable loss.
        # (That is, calculate the sum of the losses of all the training
        # set margins, and then normalize the loss by the number of
        # training examples.)
        # ===== #

        num_classes = self.W.shape[0]
        num_train = X.shape[0]

        for i in np.arange(num_train):
            ayi = self.W[y[i]].dot(X[i])
            sum_aj = 0
            for j in range(num_classes):
                sum_aj += np.exp(np.dot(self.W[j], X[i]))
```

```

        loss += (np.log(sum_aj) - ayi)
    loss /= num_train

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return loss

def loss_and_grad(self, X, y):
    """
    Same as self.loss(X, y), except that it also returns the gradient.

    Output: grad -- a matrix of the same dimensions as W containing
        the gradient of the loss with respect to W.
    """

    # Initialize the loss and gradient to zero.
    loss = 0.0
    grad = np.zeros_like(self.W)

    # ===== #
    # YOUR CODE HERE:
    # Calculate the softmax loss and the gradient. Store the gradient
    # as the variable grad.
    # ===== #
    num_classes = self.W.shape[0]
    num_train = X.shape[0]

    for i in np.arange(num_train):
        ayi = self.W[y[i]].dot(X[i])
        sum_aj = 0
        for j in range(num_classes):
            sum_aj += np.exp(np.dot(self.W[j], X[i]))
        loss += (np.log(sum_aj) - ayi)
        # calculate gradient
        for j in range(num_classes):
            aj = np.exp(np.dot(self.W[j], X[i]))
            grad[j] -= X[i] * ((j == y[i]) - aj / sum_aj)
    loss /= num_train
    grad /= num_train

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return loss, grad

def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):
    """
    sample a few random elements and only return numerical
    in these dimensions.
    """

```

```

for i in np.arange(num_checks):
    ix = tuple([np.random.randint(m) for m in self.W.shape])

    oldval = self.W[ix]
    self.W[ix] = oldval + h # increment by h
    fxph = self.loss(X, y)
    self.W[ix] = oldval - h # decrement by h
    fxmh = self.loss(X, y) # evaluate f(x - h)
    self.W[ix] = oldval # reset

    grad_numerical = (fxph - fxmh) / (2 * h)
    grad_analytic = your_grad[ix]
    rel_error = abs(grad_numerical - grad_analytic) / \
        (abs(grad_numerical) + abs(grad_analytic))
    print('numerical: %f analytic: %f, relative error: %e' %
        (grad_numerical, grad_analytic, rel_error))

def fast_loss_and_grad(self, X, y):
    """
    A vectorized implementation of loss_and_grad. It shares the same
    inputs and outputs as loss_and_grad.
    """
    loss = 0.0
    grad = np.zeros(self.W.shape) # initialize the gradient as zero

    # ===== #
    # YOUR CODE HERE:
    # Calculate the softmax loss and gradient WITHOUT any for loops.
    # ===== #
    num_train = X.shape[0]

    aj = np.dot(X, self.W.T)
    # looks like python just handles the loss fine without normalization
    # log_k = - np.max(aj, axis=1)
    ej = np.exp(aj)
    ey = ej[np.arange(num_train), y]
    sum_ej = np.sum(ej, axis=1)
    loss = -np.sum(np.log(ey/sum_ej))/num_train

    # Gradients
    probs = ej / np.sum(ej, axis=1, keepdims=True)
    probs[np.arange(num_train), y] -= 1
    grad = np.dot(probs.T, X)/num_train
    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return loss, grad

def train(self, X, y, learning_rate=1e-3, num_iters=100,
        batch_size=200, verbose=False):
    """

```

*Train this linear classifier using stochastic gradient descent.*

*Inputs:*

- *X*: A numpy array of shape  $(N, D)$  containing training data; there are  $N$  training samples each of dimension  $D$ .
- *y*: A numpy array of shape  $(N,)$  containing training labels;  $y[i] = c$  means that  $X[i]$  has label  $0 \leq c < C$  for  $C$  classes.
- *learning\_rate*: (float) learning rate for optimization.
- *num\_iters*: (integer) number of steps to take when optimizing
- *batch\_size*: (integer) number of training examples to use at each step.
- *verbose*: (boolean) If true, print progress during optimization.

*Outputs:*

*A list containing the value of the loss function at each training iteration.*

"""

```
num_train, dim = X.shape
# assume y takes values 0...K-1 where K is number of classes
num_classes = np.max(y) + 1

# initializes the weights of self.W
self.init_weights(dims=[np.max(y) + 1, X.shape[1]])

# Run stochastic gradient descent to optimize W
loss_history = []

for it in np.arange(num_iters):
    X_batch = None
    y_batch = None

    # ===== #
    # YOUR CODE HERE:
    #   Sample batch_size elements from the training data for use in
    #   gradient descent. After sampling,
    #   - X_batch should have shape: (dim, batch_size)
    #   - y_batch should have shape: (batch_size,)
    #   The indices should be randomly generated to reduce correlations
    #   in the dataset. Use np.random.choice. It's okay to sample with
    #   replacement.
    # ===== #
    idx = np.random.choice(len(X), size=batch_size, replace=False)
    X_batch = X[idx]
    y_batch = y[idx]
    # ===== #
    # END YOUR CODE HERE
    # ===== #

    # evaluate loss and gradient
    loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
    loss_history.append(loss)

    # ===== #
    # YOUR CODE HERE:
    #   Update the parameters, self.W, with a gradient step
```

```

# ===== #
self.W -= learning_rate * grad
# ===== #
# END YOUR CODE HERE
# ===== #

if verbose and it % 100 == 0:
    print('iteration {} / {}: loss {}'.format(it, num_iters, loss))

return loss_history

def predict(self, X):
    """
    Inputs:
    - X: N x D array of training data. Each row is a D-dimensional point.

    Returns:
    - y_pred: Predicted labels for the data in X. y_pred is a 1-dimensional
      array of length N, and each element is an integer giving the predicted
      class.
    """
    y_pred = np.zeros(X.shape[1])
    # ===== #
    # YOUR CODE HERE:
    # Predict the labels given the training data.
    # ===== #
    y_pred = np.argmax(np.dot(self.W, X.T), axis=0)
    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return y_pred

```