

```

1 import numpy as np
2
3 from .layers import *
4 from .layer_utils import *
5
6 """
7 This code was originally written for CS 231n at Stanford University
8 (cs231n.stanford.edu). It has been modified in various areas for use in the
9 ECE 239AS class at UCLA. This includes the descriptions of what code to
10 implement as well as some slight potential changes in variable names to be
11 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung for
12 permission to use this code. To see the original version, please visit
13 cs231n.stanford.edu.
14 """
15
16
17 class TwoLayerNet(object):
18     """
19     A two-layer fully-connected neural network with ReLU nonlinearity and
20     softmax loss that uses a modular layer design. We assume an input dimension
21     of D, a hidden dimension of H, and perform classification over C classes.
22
23     The architecture should be affine - relu - affine - softmax.
24
25     Note that this class does not implement gradient descent; instead, it
26     will interact with a separate Solver object that is responsible for running
27     optimization.
28
29     The learnable parameters of the model are stored in the dictionary
30     self.params that maps parameter names to numpy arrays.
31     """
32
33     def __init__(self, input_dim=3*32*32, hidden_dims=100, num_classes=10,
34                 dropout=0, weight_scale=1e-3, reg=0.0):
35         """
36         Initialize a new network.
37
38         Inputs:
39         - input_dim: An integer giving the size of the input
40         - hidden_dims: An integer giving the size of the hidden layer
41         - num_classes: An integer giving the number of classes to classify
42         - dropout: Scalar between 0 and 1 giving dropout strength.
43         - weight_scale: Scalar giving the standard deviation for random
44           initialization of the weights.
45         - reg: Scalar giving L2 regularization strength.
46         """
47         self.params = {}
48         self.reg = reg
49
50         # ===== #
51         # YOUR CODE HERE:
52         # Initialize W1, W2, b1, and b2. Store these as self.params['W1'],
53         # self.params['W2'], self.params['b1'] and self.params['b2']. The
54         # biases are initialized to zero and the weights are initialized
55         # so that each parameter has mean 0 and standard deviation weight_scale.
56         # The dimensions of W1 should be (input_dim, hidden_dim) and the
57         # dimensions of W2 should be (hidden_dims, num_classes)
58         # ===== #
59
60         self.params['b1'] = np.zeros(hidden_dims)

```

```

61 self.params['b2'] = np.zeros(num_classes)
62 self.params['W1'] = weight_scale * \
63     np.random.randn(input_dim, hidden_dims)
64 self.params['W2'] = weight_scale * \
65     np.random.randn(hidden_dims, num_classes)
66
67 # ===== #
68 # END YOUR CODE HERE
69 # ===== #
70
71 def loss(self, X, y=None):
72     """
73     Compute loss and gradient for a minibatch of data.
74
75     Inputs:
76     - X: Array of input data of shape (N, d_1, ..., d_k)
77     - y: Array of labels, of shape (N,). y[i] gives the label for X[i].
78
79     Returns:
80     If y is None, then run a test-time forward pass of the model and return:
81     - scores: Array of shape (N, C) giving classification scores, where
82       scores[i, c] is the classification score for X[i] and class c.
83
84     If y is not None, then run a training-time forward and backward pass and
85     return a tuple of:
86     - loss: Scalar value giving the loss
87     - grads: Dictionary with the same keys as self.params, mapping parameter
88       names to gradients of the loss with respect to those parameters.
89     """
90     scores = None
91
92     # ===== #
93     # YOUR CODE HERE:
94     # Implement the forward pass of the two-layer neural network. Store
95     # the class scores as the variable 'scores'. Be sure to use the layers
96     # you prior implemented.
97     # ===== #
98
99     h1, cache1 = affine_relu_forward(X, self.params['W1'], self.params['b1'])
100    scores, cache2 = affine_forward(h1, self.params['W2'], self.params['b2'])
101
102    # ===== #
103    # END YOUR CODE HERE
104    # ===== #
105
106    # If y is None then we are in test mode so just return scores
107    if y is None:
108        return scores
109
110    loss, grads = 0, {}
111    # ===== #
112    # YOUR CODE HERE:
113    # Implement the backward pass of the two-layer neural net. Store
114    # the loss as the variable 'loss' and store the gradients in the
115    # 'grads' dictionary. For the grads dictionary, grads['W1'] holds
116    # the gradient for W1, grads['b1'] holds the gradient for b1, etc.
117    # i.e., grads[k] holds the gradient for self.params[k].
118    #
119    # Add L2 regularization, where there is an added cost 0.5*self.reg*W^2
120    # for each W. Be sure to include the 0.5 multiplying factor to

```

```

121     # match our implementation.
122     #
123     # And be sure to use the layers you prior implemented.
124     # ===== #
125
126     loss_softmax, dscore = softmax_loss(scores, y)
127     loss_reg = self.reg * \
128         (np.sum(self.params['W1'] ** 2) +
129          np.sum(self.params['W2'] ** 2)) / 2
130     loss = loss_softmax + loss_reg
131
132     dh1, grads['W2'], grads['b2'] = affine_backward(dscore, cache2)
133     grads['W2'] += self.reg * self.params['W2']
134
135     _, grads['W1'], grads['b1'] = affine_relu_backward(dh1, cache1)
136     grads['W1'] += self.reg * self.params['W1']
137
138     # ===== #
139     # END YOUR CODE HERE
140     # ===== #
141
142     return loss, grads
143
144
145 class FullyConnectedNet(object):
146     """
147     A fully-connected neural network with an arbitrary number of hidden layers,
148     ReLU nonlinearities, and a softmax loss function. This will also implement
149     dropout and batch normalization as options. For a network with L layers,
150     the architecture will be
151
152     {affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax
153
154     where batch normalization and dropout are optional, and the {...} block is
155     repeated L - 1 times.
156
157     Similar to the TwoLayerNet above, learnable parameters are stored in the
158     self.params dictionary and will be learned using the Solver class.
159     """
160
161     def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
162                  dropout=0, use_batchnorm=False, reg=0.0,
163                  weight_scale=1e-2, dtype=np.float32, seed=None):
164         """
165         Initialize a new FullyConnectedNet.
166
167         Inputs:
168         - hidden_dims: A list of integers giving the size of each hidden layer.
169         - input_dim: An integer giving the size of the input.
170         - num_classes: An integer giving the number of classes to classify.
171         - dropout: Scalar between 0 and 1 giving dropout strength. If dropout=0 then
172           the network should not use dropout at all.
173         - use_batchnorm: Whether or not the network should use batch normalization.
174         - reg: Scalar giving L2 regularization strength.
175         - weight_scale: Scalar giving the standard deviation for random
176           initialization of the weights.
177         - dtype: A numpy datatype object; all computations will be performed using
178           this datatype. float32 is faster but less accurate, so you should use
179           float64 for numeric gradient checking.
180         - seed: If not None, then pass this random seed to the dropout layers. This

```

```

181         will make the dropout layers deterministic so we can gradient check the
182         model.
183     """
184     self.use_batchnorm = use_batchnorm
185     self.use_dropout = dropout > 0
186     self.reg = reg
187     self.num_layers = 1 + len(hidden_dims)
188     self.dtype = dtype
189     self.params = {}
190
191     # ===== #
192     # YOUR CODE HERE:
193     # Initialize all parameters of the network in the self.params dictionary.
194     # The weights and biases of layer 1 are W1 and b1; and in general the
195     # weights and biases of layer i are Wi and bi. The
196     # biases are initialized to zero and the weights are initialized
197     # so that each parameter has mean 0 and standard deviation weight_scale.
198     # ===== #
199
200     # the numbers are basically # of neurons each layer
201
202     dims = np.hstack((input_dim, hidden_dims))
203     for idx in range(1, self.num_layers):
204         self.params[('W' + str(idx))] = weight_scale * np.random.randn(dims[idx-
1], dims[idx])
205         self.params[('b' + str(idx))] = np.zeros(dims[idx])
206
207     # dims = np.array(self.hidden_dims)
208     # for idx in range(1, self.num_layers-1):
209     #     self.params[('W' + str(idx))] = weight_scale *
np.random.randn(hidden_dims[idx-1], hidden_dims[idx-1])
210     #     self.params[('b' + str(idx))] = np.zeros(hidden_dims[idx-1])
211
212
213     # ===== #
214     # END YOUR CODE HERE
215     # ===== #
216
217     # When using dropout we need to pass a dropout_param dictionary to each
218     # dropout layer so that the layer knows the dropout probability and the mode
219     # (train / test). You can pass the same dropout_param to each dropout layer.
220     self.dropout_param = {}
221     if self.use_dropout:
222         self.dropout_param = {'mode': 'train', 'p': dropout}
223         if seed is not None:
224             self.dropout_param['seed'] = seed
225
226     # With batch normalization we need to keep track of running means and
227     # variances, so we need to pass a special bn_param object to each batch
228     # normalization layer. You should pass self.bn_params[0] to the forward pass
229     # of the first batch normalization layer, self.bn_params[1] to the forward
230     # pass of the second batch normalization layer, etc.
231     self.bn_params = []
232     if self.use_batchnorm:
233         self.bn_params = [{'mode': 'train'}
for i in np.arange(self.num_layers - 1)]
234
235
236     # Cast all parameters to the correct datatype
237     for k, v in self.params.items():
238         self.params[k] = v.astype(dtype)

```

```

239
240 def loss(self, X, y=None):
241     """
242     Compute loss and gradient for the fully-connected net.
243
244     Input / output: Same as TwoLayerNet above.
245     """
246     X = X.astype(self.dtype)
247     mode = 'test' if y is None else 'train'
248
249     # Set train/test mode for batchnorm params and dropout param since they
250     # behave differently during training and testing.
251     if self.dropout_param is not None:
252         self.dropout_param['mode'] = mode
253     if self.use_batchnorm:
254         for bn_param in self.bn_params:
255             bn_param[mode] = mode
256
257     scores = None
258
259     # ===== #
260     # YOUR CODE HERE:
261     # Implement the forward pass of the FC net and store the output
262     # scores as the variable "scores".
263     # ===== #
264
265     cache = {}
266     h, cache[1] = affine_relu_forward(X, self.params['W1'], self.params['b1'])
267     for i in range(2, self.num_layers):
268         h, cache[i] = affine_relu_forward(h, self.params['W' + str(i)],
self.params['b' + str(i)])
269     scores = h
270
271     # ===== #
272     # END YOUR CODE HERE
273     # ===== #
274
275     # If test mode return early
276     if mode == 'test':
277         return scores
278
279     loss, grads = 0.0, {}
280     # ===== #
281     # YOUR CODE HERE:
282     # Implement the backwards pass of the FC net and store the gradients
283     # in the grads dict, so that grads[k] is the gradient of self.params[k]
284     # Be sure your L2 regularization includes a 0.5 factor.
285     # ===== #
286
287     loss, dscore = softmax_loss(scores, y)
288
289     W_max = 'W' + str(self.num_layers - 1)
290     b_max = 'b' + str(self.num_layers - 1)
291     dh, grads[W_max], grads[b_max] = affine_relu_backward(dscore,
cache[self.num_layers - 1])
292     for i in range(self.num_layers - 2, 0, -1):
293         dh, grads['W' + str(i)], grads['b' + str(i)] = affine_relu_backward(dh,
cache[i])
294         loss += self.reg * np.sum(self.params['W' + str(i)]**2) / 2
295         grads['W' + str(i)] += self.reg * self.params['W' + str(i)]

```

```
296  
297 # ===== #  
298 # END YOUR CODE HERE  
299 # ===== #  
300 return loss, grads  
301
```