

```

1 import numpy as np
2 import pdb
3
4 """
5 This code was originally written for CS 231n at Stanford University
6 (cs231n.stanford.edu). It has been modified in various areas for use in the
7 ECE 239AS class at UCLA. This includes the descriptions of what code to
8 implement as well as some slight potential changes in variable names to be
9 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung for
10 permission to use this code. To see the original version, please visit
11 cs231n.stanford.edu.
12 """
13
14
15 def affine_forward(x, w, b):
16     """
17     Computes the forward pass for an affine (fully-connected) layer.
18
19     The input x has shape (N, d_1, ..., d_k) and contains a minibatch of N
20     examples, where each example x[i] has shape (d_1, ..., d_k). We will
21     reshape each input into a vector of dimension D = d_1 * ... * d_k, and
22     then transform it to an output vector of dimension M.
23
24     Inputs:
25     - x: A numpy array containing input data, of shape (N, d_1, ..., d_k)
26     - w: A numpy array of weights, of shape (D, M)
27     - b: A numpy array of biases, of shape (M,)
28
29     Returns a tuple of:
30     - out: output, of shape (N, M)
31     - cache: (x, w, b)
32     """
33
34     # ===== #
35     # YOUR CODE HERE:
36     # Calculate the output of the forward pass. Notice the dimensions
37     # of w are D x M, which is the transpose of what we did in earlier
38     # assignments.
39     # ===== #
40
41     out = x.reshape(x.shape[0], -1) @ w + b
42
43
44     # ===== #
45     # END YOUR CODE HERE
46     # ===== #
47
48     cache = (x, w, b)
49     return out, cache
50
51
52 def affine_backward(dout, cache):
53     """
54     Computes the backward pass for an affine layer.
55
56     Inputs:
57     - dout: Upstream derivative, of shape (N, M)
58     - cache: Tuple of:
59         - x: Input data, of shape (N, d_1, ... d_k)
60         - w: Weights, of shape (D, M)

```

```

61
62 Returns a tuple of:
63 - dx: Gradient with respect to x, of shape (N, d1, ..., d_k)
64 - dw: Gradient with respect to w, of shape (D, M)
65 - db: Gradient with respect to b, of shape (M,)
66 """
67 x, w, b = cache
68 dx, dw, db = None, None, None
69
70 # ===== #
71 # YOUR CODE HERE:
72 # Calculate the gradients for the backward pass.
73 # ===== #
74
75 # dout is N x M
76 # dx should be N x d1 x ... x dk; it relates to dout through multiplication with
w, which is D x M
77 # dw should be D x M; it relates to dout through multiplication with x, which is
N x D after reshaping
78 # db should be M; it is just the sum over dout examples
79
80 x_ND = x.reshape(x.shape[0], -1) #(N, D)
81 dx = (dout @ w.T).reshape(x.shape)
82 dw = x_ND.T @ dout
83 db = np.sum(dout, axis=0)
84
85
86 # ===== #
87 # END YOUR CODE HERE
88 # ===== #
89
90 return dx, dw, db
91
92
93 def relu_forward(x):
94     """
95     Computes the forward pass for a layer of rectified linear units (ReLU).
96
97     Input:
98     - x: Inputs, of any shape
99
100     Returns a tuple of:
101     - out: Output, of the same shape as x
102     - cache: x
103     """
104     # ===== #
105     # YOUR CODE HERE:
106     # Implement the ReLU forward pass.
107     # ===== #
108
109     relu = lambda x: x * (x > 0)
110     out = relu(x)
111     # ===== #
112     # END YOUR CODE HERE
113     # ===== #
114
115     cache = x
116     return out, cache
117
118

```

```

119 def relu_backward(dout, cache):
120     """
121     Computes the backward pass for a layer of rectified linear units (ReLU).
122
123     Input:
124     - dout: Upstream derivatives, of any shape
125     - cache: Input x, of same shape as dout
126
127     Returns:
128     - dx: Gradient with respect to x
129     """
130     x = cache
131
132     # ===== #
133     # YOUR CODE HERE:
134     #   Implement the ReLU backward pass
135     # ===== #
136
137     # ReLU directs linearly to those > 0
138     dx = dout * (x > 0)
139
140     # ===== #
141     # END YOUR CODE HERE
142     # ===== #
143
144     return dx
145
146 def svm_loss(x, y):
147     """
148     Computes the loss and gradient using for multiclass SVM classification.
149
150     Inputs:
151     - x: Input data, of shape (N, C) where x[i, j] is the score for the jth class
152         for the ith input.
153     - y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and
154         0 <= y[i] < C
155
156     Returns a tuple of:
157     - loss: Scalar giving the loss
158     - dx: Gradient of the loss with respect to x
159     """
160
161     N = x.shape[0]
162     correct_class_scores = x[np.arange(N), y]
163     margins = np.maximum(0, x - correct_class_scores[:, np.newaxis] + 1.0)
164     margins[np.arange(N), y] = 0
165     loss = np.sum(margins) / N
166     num_pos = np.sum(margins > 0, axis=1)
167     dx = np.zeros_like(x)
168     dx[margins > 0] = 1
169     dx[np.arange(N), y] -= num_pos
170     dx /= N
171     return loss, dx
172
173 def softmax_loss(x, y):
174     """
175     Computes the loss and gradient for softmax classification.
176
177     Inputs:

```

```
179 - x: Input data, of shape (N, C) where x[i, j] is the score for the jth class
180     for the ith input.
181 - y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and
182     0 <= y[i] < C
183
184 Returns a tuple of:
185 - loss: Scalar giving the loss
186 - dx: Gradient of the loss with respect to x
187 """
188
189 probs = np.exp(x - np.max(x, axis=1, keepdims=True))
190 probs /= np.sum(probs, axis=1, keepdims=True)
191 N = x.shape[0]
192 loss = -np.sum(np.log(probs[np.arange(N), y])) / N
193 dx = probs.copy()
194 dx[np.arange(N), y] -= 1
195 dx /= N
196 return loss, dx
197
```