

```

1 from os import replace
2 import numpy as np
3 import matplotlib.pyplot as plt
4
5 """
6 This code was originally written for CS 231n at Stanford University
7 (cs231n.stanford.edu). It has been modified in various areas for use in the
8 ECE 239AS class at UCLA. This includes the descriptions of what code to
9 implement as well as some slight potential changes in variable names to be
10 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung for
11 permission to use this code. To see the original version, please visit
12 cs231n.stanford.edu.
13 """
14
15
16 class TwoLayerNet(object):
17     """
18     A two-layer fully-connected neural network. The net has an input dimension of
19     N, a hidden layer dimension of H, and performs classification over C classes.
20     We train the network with a softmax loss function and L2 regularization on the
21     weight matrices. The network uses a ReLU nonlinearity after the first fully
22     connected layer.
23
24     In other words, the network has the following architecture:
25
26     input - fully connected layer - ReLU - fully connected layer - softmax
27
28     The outputs of the second fully-connected layer are the scores for each class.
29     """
30
31     def __init__(self, input_size, hidden_size, output_size, std=1e-4):
32         """
33         Initialize the model. Weights are initialized to small random values and
34         biases are initialized to zero. Weights and biases are stored in the
35         variable self.params, which is a dictionary with the following keys:
36
37         W1: First layer weights; has shape (H, D)
38         b1: First layer biases; has shape (H,)
39         W2: Second layer weights; has shape (C, H)
40         b2: Second layer biases; has shape (C,)
41
42         Inputs:
43         - input_size: The dimension D of the input data.
44         - hidden_size: The number of neurons H in the hidden layer.
45         - output_size: The number of classes C.
46         """
47         self.params = {}
48         self.params['W1'] = std * np.random.randn(hidden_size, input_size)
49         self.params['b1'] = np.zeros(hidden_size)
50         self.params['W2'] = std * np.random.randn(output_size, hidden_size)
51         self.params['b2'] = np.zeros(output_size)
52
53     def loss(self, X, y=None, reg=0.0):
54         """
55         Compute the loss and gradients for a two layer fully connected neural
56         network.
57
58         Inputs:
59         - X: Input data of shape (N, D). Each X[i] is a training sample.
60         - y: Vector of training labels. y[i] is the label for X[i], and each y[i] is

```

```

61         an integer in the range  $0 \leq y[i] < C$ . This parameter is optional; if it
62         is not passed then we only return scores, and if it is passed then we
63         instead return the loss and gradients.
64     - reg: Regularization strength.
65
66     Returns:
67     If y is None, return a matrix scores of shape (N, C) where scores[i, c] is
68     the score for class c on input X[i].
69
70     If y is not None, instead return a tuple of:
71     - loss: Loss (data loss and regularization loss) for this batch of training
72       samples.
73     - grads: Dictionary mapping parameter names to gradients of those parameters
74       with respect to the loss function; has the same keys as self.params.
75     """
76     # Unpack variables from the params dictionary
77     W1, b1 = self.params['W1'], self.params['b1']
78     W2, b2 = self.params['W2'], self.params['b2']
79     N, D = X.shape
80
81     # Compute the forward pass
82     scores = None
83
84     # ===== #
85     # YOUR CODE HERE:
86     # Calculate the output scores of the neural network. The result
87     # should be (N, C). As stated in the description for this class,
88     # there should not be a ReLU layer after the second FC layer.
89     # The output of the second FC layer is the output scores. Do not
90     # use a for loop in your implementation.
91     # ===== #
92
93     relu = lambda x: x*(x > 0)
94     h1 = relu(X @ W1.T + b1)
95     scores = h1 @ W2.T + b2
96
97     # ===== #
98     # END YOUR CODE HERE
99     # ===== #
100
101     # If the targets are not given then jump out, we're done
102     if y is None:
103         return scores
104
105     # Compute the loss
106     loss = None
107
108     # ===== #
109     # YOUR CODE HERE:
110     # Calculate the loss of the neural network. This includes the
111     # softmax loss and the L2 regularization for W1 and W2. Store the
112     # total loss in the variable loss. Multiply the regularization
113     # loss by 0.5 (in addition to the factor reg).
114     # ===== #
115
116     # scores is num_examples by num_classes
117
118     exp_scores = np.exp(scores)
119     probs = exp_scores / np.sum(exp_scores, axis=1, keepdims=True)
120     true_probs = -np.log(probs[np.arange(N), y])

```

```

121 data_loss = np.sum(true_probs) / N
122
123 reg_loss = (np.sum(W1**2) + np.sum(W2**2))*reg/2
124 loss = data_loss + reg_loss
125
126 # ===== #
127 # END YOUR CODE HERE
128 # ===== #
129
130 grads = {}
131
132 # ===== #
133 # YOUR CODE HERE:
134 # Implement the backward pass. Compute the derivatives of the
135 # weights and the biases. Store the results in the grads
136 # dictionary. e.g., grads['W1'] should store the gradient for
137 # W1, and be of the same size as W1.
138 # ===== #
139
140 dz2 = probs
141 dz2[np.arange(N), y] -= 1
142 dz2 /= N
143
144 grads["W2"] = dz2.T @ h1 + reg * W2
145 grads["b2"] = dz2.T @ np.ones(N)
146
147 dz1 = dz2 @ W2 * (h1 > 0)
148
149 grads["W1"] = dz1.T @ X + reg * W1
150 grads["b1"] = dz1.T @ np.ones(N)
151
152 # ===== #
153 # END YOUR CODE HERE
154 # ===== #
155
156 return loss, grads
157
158 def train(self, X, y, X_val, y_val,
159           learning_rate=1e-3, learning_rate_decay=0.95,
160           reg=1e-5, num_iters=100,
161           batch_size=200, verbose=False):
162     """
163     Train this neural network using stochastic gradient descent.
164
165     Inputs:
166     - X: A numpy array of shape (N, D) giving training data.
167     - y: A numpy array of shape (N,) giving training labels; y[i] = c means that
168         X[i] has label c, where 0 <= c < C.
169     - X_val: A numpy array of shape (N_val, D) giving validation data.
170     - y_val: A numpy array of shape (N_val,) giving validation labels.
171     - learning_rate: Scalar giving learning rate for optimization.
172     - learning_rate_decay: Scalar giving factor used to decay the learning rate
173         after each epoch.
174     - reg: Scalar giving regularization strength.
175     - num_iters: Number of steps to take when optimizing.
176     - batch_size: Number of training examples to use per step.
177     - verbose: boolean; if true print progress during optimization.
178     """
179     num_train = X.shape[0]
180     # iterations_per_epoch = max(num_train / batch_size, 1)

```

```

181 iterations_per_epoch = max(int(num_train / batch_size), 1)
182
183 # Use SGD to optimize the parameters in self.model
184 loss_history = []
185 train_acc_history = []
186 val_acc_history = []
187
188 for it in np.arange(num_iters):
189     X_batch = None
190     y_batch = None
191
192     # ===== #
193     # YOUR CODE HERE:
194     #   Create a minibatch by sampling batch_size samples randomly.
195     # ===== #
196     idx = np.random.choice(len(X), size=batch_size)
197     X_batch = X[idx]
198     y_batch = y[idx]
199     # ===== #
200     # END YOUR CODE HERE
201     # ===== #
202
203     # Compute loss and gradients using the current minibatch
204     loss, grads = self.loss(X_batch, y=y_batch, reg=reg)
205     loss_history.append(loss)
206
207     # ===== #
208     # YOUR CODE HERE:
209     #   Perform a gradient descent step using the minibatch to update
210     #   all parameters (i.e., W1, W2, b1, and b2).
211     # ===== #
212
213     self.params['W1'] -= learning_rate * grads['W1']
214     self.params['b1'] -= learning_rate * grads['b1']
215     self.params['W2'] -= learning_rate * grads['W2']
216     self.params['b2'] -= learning_rate * grads['b2']
217
218     # ===== #
219     # END YOUR CODE HERE
220     # ===== #
221
222     if verbose and it % 100 == 0:
223         print('iteration {} / {}: loss {}'.format(it, num_iters, loss))
224
225     # Every epoch, check train and val accuracy and decay learning rate.
226     if it % iterations_per_epoch == 0:
227         # Check accuracy
228         train_acc = (self.predict(X_batch) == y_batch).mean()
229         val_acc = (self.predict(X_val) == y_val).mean()
230         train_acc_history.append(train_acc)
231         val_acc_history.append(val_acc)
232
233         # Decay learning rate
234         learning_rate *= learning_rate_decay
235
236     return {
237         'loss_history': loss_history,
238         'train_acc_history': train_acc_history,
239         'val_acc_history': val_acc_history,
240     }

```

```
241
242 def predict(self, X):
243     """
244     Use the trained weights of this two-layer network to predict labels for
245     data points. For each data point we predict scores for each of the C
246     classes, and assign each data point to the class with the highest score.
247
248     Inputs:
249     - X: A numpy array of shape (N, D) giving N D-dimensional data points to
250         classify.
251
252     Returns:
253     - y_pred: A numpy array of shape (N,) giving predicted labels for each of
254         the elements of X. For all i, y_pred[i] = c means that X[i] is predicted
255         to have class c, where 0 ≤ c < C.
256     """
257     y_pred = None
258
259     # ===== #
260     # YOUR CODE HERE:
261     #   Predict the class given the input data.
262     # ===== #
263
264     y_pred = np.argmax(self.loss(X), axis=1)
265
266     # ===== #
267     # END YOUR CODE HERE
268     # ===== #
269
270     return y_pred
271
```