

# two\_layer\_nn

February 3, 2021

## 0.1 This is the 2-layer neural network workbook for ECE 247 Assignment #3

Please follow the notebook linearly to implement a two layer neural network.

Please print out the workbook entirely when completed.

We thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class (cs231n.stanford.edu). These are the functions in the cs231n folders and code in the jupyter notebook to preprocess and show the images. The classifiers used are based off of code prepared for CS 231n as well.

The goal of this workbook is to give you experience with training a two layer neural network.

```
[1]: import random
import numpy as np
from cs231n.data_utils import load_CIFAR10
import matplotlib.pyplot as plt

%matplotlib inline
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

## 0.2 Toy example

Before loading CIFAR-10, there will be a toy example to test your implementation of the forward and backward pass

```
[2]: from nndl.neural_net import TwoLayerNet

[3]: # Create a small net and some toy data to check your implementations.
# Note that we set the random seed for repeatable experiments.

input_size = 4
hidden_size = 10
num_classes = 3
num_inputs = 5
```

```

def init_toy_model():
    np.random.seed(0)
    return TwoLayerNet(input_size, hidden_size, num_classes, std=1e-1)

def init_toy_data():
    np.random.seed(1)
    X = 10 * np.random.randn(num_inputs, input_size)
    y = np.array([0, 1, 2, 2, 1])
    return X, y

net = init_toy_model()
X, y = init_toy_data()

```

### 0.2.1 Compute forward pass scores

```

[4]: ## Implement the forward pass of the neural network.

# Note, there is a statement if y is None: return scores, which is why
# the following call will calculate the scores.
scores = net.loss(X)
print('Your scores:')
print(scores)
print()
print('correct scores:')
correct_scores = np.asarray([
    [-1.07260209,  0.05083871, -0.87253915],
    [-2.02778743, -0.10832494, -1.52641362],
    [-0.74225908,  0.15259725, -0.39578548],
    [-0.38172726,  0.10835902, -0.17328274],
    [-0.64417314, -0.18886813, -0.41106892]])
print(correct_scores)
print()

# The difference should be very small. We get < 1e-7
print('Difference between your scores and correct scores:')
print(np.sum(np.abs(scores - correct_scores)))

```

Your scores:

```

[[-1.07260209  0.05083871 -0.87253915]
 [-2.02778743 -0.10832494 -1.52641362]
 [-0.74225908  0.15259725 -0.39578548]
 [-0.38172726  0.10835902 -0.17328274]
 [-0.64417314 -0.18886813 -0.41106892]]

```

correct scores:

```

[[-1.07260209  0.05083871 -0.87253915]
 [-2.02778743 -0.10832494 -1.52641362]]

```

```
[-0.74225908  0.15259725 -0.39578548]
[-0.38172726  0.10835902 -0.17328274]
[-0.64417314 -0.18886813 -0.41106892]]
```

Difference between your scores and correct scores:  
3.381231197113754e-08

### 0.2.2 Forward pass loss

```
[5]: loss, _ = net.loss(X, y, reg=0.05)
correct_loss = 1.071696123862817

# should be very small, we get < 1e-12
print('Difference between your loss and correct loss:')
print(np.sum(np.abs(loss - correct_loss)))
```

Difference between your loss and correct loss:  
0.0

```
[6]: print(loss)
```

1.071696123862817

### 0.2.3 Backward pass

Implements the backwards pass of the neural network. Check your gradients with the gradient check utilities provided.

```
[7]: from cs231n.gradient_check import eval_numerical_gradient

# Use numeric gradient checking to check your implementation of the backward
    ↪pass.
# If your implementation is correct, the difference between the numeric and
# analytic gradients should be less than 1e-8 for each of W1, W2, b1, and b2.

loss, grads = net.loss(X, y, reg=0.05)

# these should all be less than 1e-8 or so
for param_name in grads:
    f = lambda W: net.loss(X, y, reg=0.05)[0]
    param_grad_num = eval_numerical_gradient(f, net.params[param_name],
    ↪verbose=False)
    print('{} max relative error: {}'.format(param_name,
    ↪rel_error(param_grad_num, grads[param_name])))
```

W2 max relative error: 2.9632221903873815e-10  
b2 max relative error: 1.2482624742512528e-09  
W1 max relative error: 1.283285096965795e-09  
b1 max relative error: 3.172680285697327e-09

### 0.2.4 Training the network


Implement `neural_net.train()` to train the network via stochastic gradient descent, much like the softmax and SVM.

```
[8]: net = init_toy_model()
stats = net.train(X, y, X, y,
                  learning_rate=1e-1, reg=5e-6,
                  num_iters=100, verbose=False)

print('Final training loss: ', stats['loss_history'][-1])

# plot the loss history
plt.plot(stats['loss_history'])
plt.xlabel('iteration')
plt.ylabel('training loss')
plt.title('Training Loss history')
plt.show()
```

Final training loss: 0.014497864587765906



## 0.3 Classify CIFAR-10

Do classification on the CIFAR-10 dataset.

```
[9]: from cs231n.data_utils import load_CIFAR10

def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000):
    """
    Load the CIFAR-10 dataset from disk and perform preprocessing to prepare
    it for the two-layer neural net classifier. These are the same steps as
    we used for the SVM, but condensed to a single function.
    """
    # Load the raw CIFAR-10 data
    cifar10_dir = '../cifar-10-batches-py'
    X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

    # Subsample the data
    mask = list(range(num_training, num_training + num_validation))
    X_val = X_train[mask]
    y_val = y_train[mask]
    mask = list(range(num_training))
    X_train = X_train[mask]
    y_train = y_train[mask]
    mask = list(range(num_test))
    X_test = X_test[mask]
    y_test = y_test[mask]

    # Normalize the data: subtract the mean image
    mean_image = np.mean(X_train, axis=0)
    X_train -= mean_image
    X_val -= mean_image
    X_test -= mean_image

    # Reshape data to rows
    X_train = X_train.reshape(num_training, -1)
    X_val = X_val.reshape(num_validation, -1)
    X_test = X_test.reshape(num_test, -1)

    return X_train, y_train, X_val, y_val, X_test, y_test

# Invoke the above function to get our data.
X_train, y_train, X_val, y_val, X_test, y_test = get_CIFAR10_data()
print('Train data shape: ', X_train.shape)
print('Train labels shape: ', y_train.shape)
print('Validation data shape: ', X_val.shape)
print('Validation labels shape: ', y_val.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
```

Train data shape: (49000, 3072)

Train labels shape: (49000,)

```
Validation data shape: (1000, 3072)
Validation labels shape: (1000,)
Test data shape: (1000, 3072)
Test labels shape: (1000,)
```

### 0.3.1 Running SGD

If your implementation is correct, you should see a validation accuracy of around 28-29%.

```
[10]: input_size = 32 * 32 * 3
      hidden_size = 50
      num_classes = 10
      net = TwoLayerNet(input_size, hidden_size, num_classes)

      # Train the network
      stats = net.train(X_train, y_train, X_val, y_val,
                        num_iters=1000, batch_size=200,
                        learning_rate=1e-4, learning_rate_decay=0.95,
                        reg=0.25, verbose=True)

      # Predict on the validation set
      val_acc = (net.predict(X_val) == y_val).mean()
      print('Validation accuracy: ', val_acc)

      # Save this net as the variable subopt_net for later comparison.
      subopt_net = net
```

```
iteration 0 / 1000: loss 2.302757518613176
iteration 100 / 1000: loss 2.302120159207236
iteration 200 / 1000: loss 2.2956136007408703
iteration 300 / 1000: loss 2.2518259043164135
iteration 400 / 1000: loss 2.188995235046776
iteration 500 / 1000: loss 2.1162527791897747
iteration 600 / 1000: loss 2.064670827698217
iteration 700 / 1000: loss 1.9901688623083942
iteration 800 / 1000: loss 2.002827640124685
iteration 900 / 1000: loss 1.9465176817856495
Validation accuracy: 0.283
```

## 0.4 Questions:


The training accuracy isn't great.

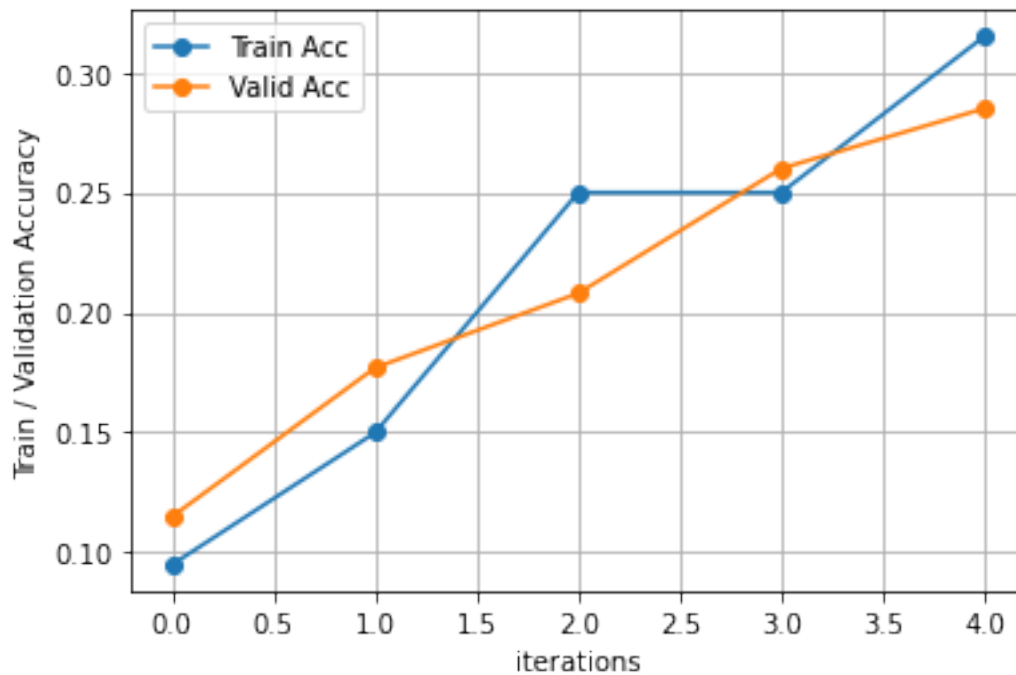
- (1) What are some of the reasons why this is the case? Take the following cell to do some analyses and then report your answers in the cell following the one below.
- (2) How should you fix the problems you identified in (1)?

```
[11]: stats['train_acc_history']
```

```
[11]: [0.095, 0.15, 0.25, 0.25, 0.315]
```

```
[12]: # ===== #  
# YOUR CODE HERE:  
# Do some debugging to gain some insight into why the optimization  
# isn't great.  
# ===== #  
  
# Plot the loss function and train / validation accuracies  
  
plt.plot(stats['loss_history'])  
plt.xlabel('iterations')  
plt.ylabel('Loss')  
plt.grid()  
plt.show()  
  
plt.plot(stats['train_acc_history'], marker='o', label='Train Acc')  
plt.plot(stats['val_acc_history'], marker='o', label="Valid Acc")  
plt.xlabel('iterations')  
plt.ylabel('Train / Validation Accuracy')  
plt.legend()  
plt.grid()  
plt.show()  
# ===== #  
# END YOUR CODE HERE  
# ===== #
```





## 0.5 Answers:

- (1) As shown in the Iteration vs Loss plot, the loss started to fluctuate as iteration number increases, which can be caused by overshooting under the current learning rate, or it can be an inherent zigzagging issue with ReLU. Also, the training and validation accuracy haven't plateaued before the training ends, which means the model is probably underfitted.
- (2) Decreasing the learning rate and increasing batch size can help with the fluctuating loss; increasing the number of training epochs can help with underfitting.

## 0.6 Optimize the neural network

Use the following part of the Jupyter notebook to optimize your hyperparameters on the validation set. Store your nets as `best_net`.

```
[13]: best_net = None # store the best model into this

# ===== #
# YOUR CODE HERE:
# Optimize over your hyperparameters to arrive at the best neural
# network. You should be able to get over 50% validation accuracy.
# For this part of the notebook, we will give credit based on the
# accuracy you get. Your score on this question will be multiplied by:
# min(floor((X - 28%)) / %22, 1)
```



```

# where if you get 50% or higher validation accuracy, you get full
# points.
#
# Note, you need to use the same network structure (keep hidden_size = 50)!
# ===== #

best_lr = 1e-3
best_decay = 0.95
best_reg = 1e-5
best_num_iters = 100
best_size = 200

lrs = []
y_val_accs = []

print("testing learning rates:")
for lr in range(10):
    lr = 0.01*10**(-lr)
    lrs.append(lr)
    net = TwoLayerNet(input_size, hidden_size, num_classes)
    net.train(X_train, y_train, X_val, y_val, learning_rate=lr)
    y_pred = net.predict(X_val)
    accuracy = 1 - np.count_nonzero(y_val - y_pred) / y_val.shape[0]
    y_val_accs.append(accuracy)

    # print("\tlearning rate = ", lr, "; Accuracy = ", accuracy)
best_idx = np.argmax(y_val_accs)
best_lr = lrs[best_idx]
print("\tBest learning rate =", best_lr, "Val acc =", y_val_accs[best_idx])


plt.plot(lrs, y_val_accs)
plt.show()

```

```

testing learning rates:
/w/home.13/class/classshan/Homework/HW3/nndl/neural_net.py:120: RuntimeWarning:
divide by zero encountered in log
    true_probs = -np.log(probs[np.arange(N), y])
/w/home.13/class/classshan/Homework/HW3/nndl/neural_net.py:118: RuntimeWarning:
overflow encountered in exp
    exp_scores = np.exp(scores)
/w/home.13/class/classshan/Homework/HW3/nndl/neural_net.py:119: RuntimeWarning:
invalid value encountered in true_divide
    probs = exp_scores / np.sum(exp_scores, axis=1, keepdims=True)
    Best learning rate = 0.001 Val acc = 0.281

```




```
[14]: print("testing num_iters:")
num_iters = 100*np.arange(1, 21)
y_val_accs = []

for num_iter in num_iters:
    net = TwoLayerNet(input_size, hidden_size, num_classes)
    net.train(X_train, y_train, X_val, y_val, learning_rate=best_lr,
    ↪ num_iters=num_iter)
    y_pred = net.predict(X_val)
    accuracy = 1 - np.count_nonzero(y_val - y_pred) / y_val.shape[0]
    if accuracy > 0.5:
        best_net = net
    y_val_accs.append(accuracy)

best_idx = np.argmax(y_val_accs)
best_num_iters = num_iters[best_idx]
print("\tBest number of iterations =", best_num_iters, "Val acc =",
    ↪ y_val_accs[best_idx])

plt.plot(num_iters, y_val_accs)
plt.show()
```

```
testing num_iters:
    Best number of iterations = 1800 Val acc = 0.511
```



```
[15]: print("testing regs:")
regs = []
y_val_accs = []

for reg in range(3, 8):
    reg = 10**(-reg)
    regs.append(reg)
    net = TwoLayerNet(input_size, hidden_size, num_classes)
    net.train(X_train, y_train, X_val, y_val,
              learning_rate=best_lr,
              num_iters=best_num_iters,
              reg=reg)
    y_pred = net.predict(X_val)
    accuracy = 1 - np.count_nonzero(y_val - y_pred) / y_val.shape[0]
    y_val_accs.append(accuracy)

best_idx = np.argmax(y_val_accs)
best_reg = regs[best_idx]
print("\tBest reg =", best_reg, "Val acc =", y_val_accs[best_idx])

print("testing batch size:")
batch_sizes = 50 * np.arange(1, 11)
y_val_accs = []

for size in batch_sizes:
```

```

net = TwoLayerNet(input_size, hidden_size, num_classes)
net.train(X_train, y_train, X_val, y_val,
          learning_rate=best_lr,
          num_iters=best_num_iters,
          reg=best_reg,
          batch_size=size)
y_pred = net.predict(X_val)
accuracy = 1 - np.count_nonzero(y_val - y_pred) / y_val.shape[0]
y_val_accs.append(accuracy)

best_idx = np.argmax(y_val_accs)
best_size = batch_sizes[best_idx]
print("\tBest batch size =", best_size, "Val acc =", y_val_accs[best_idx])

print("testing lr decay:")
lr_decays = []
y_val_accs = []

for decay in range(9):
    decay = 1 - 0.025*decay
    lr_decays.append(decay)
    net = TwoLayerNet(input_size, hidden_size, num_classes)
    net.train(X_train, y_train, X_val, y_val,
              learning_rate=best_lr,
              num_iters=best_num_iters,
              reg=best_reg,
              batch_size=best_size,
              learning_rate_decay=decay)
    y_pred = net.predict(X_val)
    accuracy = 1 - np.count_nonzero(y_val - y_pred) / y_val.shape[0]
    y_val_accs.append(accuracy)
best_idx = np.argmax(y_val_accs)
best_decay = lr_decays[best_idx]
print("\tBest lr decay =", best_decay, "Val acc =", y_val_accs[best_idx])

```

testing regs:

Best reg = 1e-05 Val acc = 0.498

testing batch size:

Best batch size = 500 Val acc = 0.5

testing lr decay:

Best lr decay = 0.975 Val acc = 0.517

[37]: *# Training Best Net using optimized parameters*

```

net = TwoLayerNet(input_size, hidden_size, num_classes)
net.train(X_train, y_train, X_val, y_val,
          learning_rate=best_lr,

```

```

        num_iters=best_num_iters,
        reg=best_reg,
        batch_size=best_size,
        learning_rate_decay=best_decay)
y_pred = net.predict(X_val)
accuracy = 1 - np.count_nonzero(y_val - y_pred) / y_val.shape[0]
print("validation accuracy =", accuracy)

# ===== #
# END YOUR CODE HERE
# ===== #
best_net = net

```

validation accuracy = 0.515

```


[39]: from cs231n.vis_utils import visualize_grid


# Visualize the weights of the network

def show_net_weights(net):
    W1 = net.params['W1']
    W1 = W1.T.reshape(32, 32, 3, -1).transpose(3, 0, 1, 2)
    plt.imshow(visualize_grid(W1, padding=3).astype('uint8'))
    plt.gca().axis('off')
    plt.show()

show_net_weights(subopt_net)
show_net_weights(best_net)

```





### 0.7 Question:

- (1) What differences do you see in the weights between the suboptimal net and the best net you arrived at?

### 0.8 Answer:

- (1) The weights in the suboptimal net has a lot of similarities between them, both in terms of color and shape. They also have simple shape and are not very evolved. The weights of the best net are highly evolved with high resolutions and have little similarities between them.

### 0.9 Evaluate on test set

```
[38]: test_acc = (net.predict(X_test) == y_test).mean()
      print('Test accuracy: ', test_acc)
```

Test accuracy: 0.504