

```

1 import numpy as np
2 import pdb
3
4 from .layers import *
5 from .layer_utils import *
6
7 """
8 This code was originally written for CS 231n at Stanford University
9 (cs231n.stanford.edu). It has been modified in various areas for use in the
10 ECE 239AS class at UCLA. This includes the descriptions of what code to
11 implement as well as some slight potential changes in variable names to be
12 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung for
13 permission to use this code. To see the original version, please visit
14 cs231n.stanford.edu.
15 """
16
17
18 class TwoLayerNet(object):
19     """
20     A two-layer fully-connected neural network with ReLU nonlinearity and
21     softmax loss that uses a modular layer design. We assume an input dimension
22     of D, a hidden dimension of H, and perform classification over C classes.
23
24     The architecture should be affine - relu - affine - softmax.
25
26     Note that this class does not implement gradient descent; instead, it
27     will interact with a separate Solver object that is responsible for running
28     optimization.
29
30     The learnable parameters of the model are stored in the dictionary
31     self.params that maps parameter names to numpy arrays.
32     """
33
34     def __init__(self, input_dim=3*32*32, hidden_dims=100, num_classes=10,
35                 dropout=0, weight_scale=1e-3, reg=0.0):
36         """
37         Initialize a new network.
38
39         Inputs:
40         - input_dim: An integer giving the size of the input
41         - hidden_dims: An integer giving the size of the hidden layer
42         - num_classes: An integer giving the number of classes to classify
43         - dropout: Scalar between 0 and 1 giving dropout strength.
44         - weight_scale: Scalar giving the standard deviation for random
45           initialization of the weights.
46         - reg: Scalar giving L2 regularization strength.
47         """
48         self.params = {}
49         self.reg = reg
50
51         # ===== #
52         # YOUR CODE HERE:
53         # Initialize W1, W2, b1, and b2. Store these as self.params['W1'],
54         # self.params['W2'], self.params['b1'] and self.params['b2']. The
55         # biases are initialized to zero and the weights are initialized
56         # so that each parameter has mean 0 and standard deviation weight_scale.
57         # The dimensions of W1 should be (input_dim, hidden_dim) and the
58         # dimensions of W2 should be (hidden_dims, num_classes)
59         # ===== #
60         self.params['b1'] = np.zeros(hidden_dims)

```

```

61 self.params['b2'] = np.zeros(num_classes)
62 self.params['W1'] = weight_scale * \
63     np.random.randn(input_dim, hidden_dims)
64 self.params['W2'] = weight_scale * \
65     np.random.randn(hidden_dims, num_classes)
66 # ===== #
67 # END YOUR CODE HERE
68 # ===== #
69
70 def loss(self, X, y=None):
71     """
72     Compute loss and gradient for a minibatch of data.
73
74     Inputs:
75     - X: Array of input data of shape (N, d_1, ..., d_k)
76     - y: Array of labels, of shape (N,). y[i] gives the label for X[i].
77
78     Returns:
79     If y is None, then run a test-time forward pass of the model and return:
80     - scores: Array of shape (N, C) giving classification scores, where
81       scores[i, c] is the classification score for X[i] and class c.
82
83     If y is not None, then run a training-time forward and backward pass and
84     return a tuple of:
85     - loss: Scalar value giving the loss
86     - grads: Dictionary with the same keys as self.params, mapping parameter
87       names to gradients of the loss with respect to those parameters.
88     """
89     scores = None
90
91     # ===== #
92     # YOUR CODE HERE:
93     # Implement the forward pass of the two-layer neural network. Store
94     # the class scores as the variable 'scores'. Be sure to use the layers
95     # you prior implemented.
96     # ===== #
97     h1, cache1 = affine_relu_forward(
98         X, self.params['W1'], self.params['b1'])
99     scores, cache2 = affine_forward(
100         h1, self.params['W2'], self.params['b2'])
101     # ===== #
102     # END YOUR CODE HERE
103     # ===== #
104
105     # If y is None then we are in test mode so just return scores
106     if y is None:
107         return scores
108
109     loss, grads = 0, {}
110     # ===== #
111     # YOUR CODE HERE:
112     # Implement the backward pass of the two-layer neural net. Store
113     # the loss as the variable 'loss' and store the gradients in the
114     # 'grads' dictionary. For the grads dictionary, grads['W1'] holds
115     # the gradient for W1, grads['b1'] holds the gradient for b1, etc.
116     # i.e., grads[k] holds the gradient for self.params[k].
117     #
118     # Add L2 regularization, where there is an added cost 0.5*self.reg*W^2
119     # for each W. Be sure to include the 0.5 multiplying factor to
120     # match our implementation.

```

```

121     #
122     # And be sure to use the layers you prior implemented.
123     # ===== #
124     loss_softmax, dscore = softmax_loss(scores, y)
125     loss_reg = self.reg * \
126         (np.sum(self.params['W1'] ** 2) +
127          np.sum(self.params['W2'] ** 2)) / 2
128     loss = loss_softmax + loss_reg
129
130     dh1, grads['W2'], grads['b2'] = affine_backward(dscore, cache2)
131     grads['W2'] += self.reg * self.params['W2']
132
133     _, grads['W1'], grads['b1'] = affine_relu_backward(dh1, cache1)
134     grads['W1'] += self.reg * self.params['W1']
135     # ===== #
136     # END YOUR CODE HERE
137     # ===== #
138
139     return loss, grads
140
141
142 class FullyConnectedNet(object):
143     """
144     A fully-connected neural network with an arbitrary number of hidden layers,
145     ReLU nonlinearities, and a softmax loss function. This will also implement
146     dropout and batch normalization as options. For a network with L layers,
147     the architecture will be
148
149     {affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax
150
151     where batch normalization and dropout are optional, and the {...} block is
152     repeated L - 1 times.
153
154     Similar to the TwoLayerNet above, learnable parameters are stored in the
155     self.params dictionary and will be learned using the Solver class.
156     """
157
158     def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
159                 dropout=0, use_batchnorm=False, reg=0.0,
160                 weight_scale=1e-2, dtype=np.float32, seed=None):
161         """
162         Initialize a new FullyConnectedNet.
163
164         Inputs:
165         - hidden_dims: A list of integers giving the size of each hidden layer.
166         - input_dim: An integer giving the size of the input.
167         - num_classes: An integer giving the number of classes to classify.
168         - dropout: Scalar between 0 and 1 giving dropout strength. If dropout=0 then
169           the network should not use dropout at all.
170         - use_batchnorm: Whether or not the network should use batch normalization.
171         - reg: Scalar giving L2 regularization strength.
172         - weight_scale: Scalar giving the standard deviation for random
173           initialization of the weights.
174         - dtype: A numpy datatype object; all computations will be performed using
175           this datatype. float32 is faster but less accurate, so you should use
176           float64 for numeric gradient checking.
177         - seed: If not None, then pass this random seed to the dropout layers. This
178           will make the dropout layers deterministic so we can gradient check the
179           model.
180         """

```

```

181 self.use_batchnorm = use_batchnorm
182 self.use_dropout = dropout > 0
183 self.reg = reg
184 self.num_layers = 1 + len(hidden_dims)
185 self.dtype = dtype
186 self.params = {}
187
188 # ===== #
189 # YOUR CODE HERE:
190 # Initialize all parameters of the network in the self.params dictionary.
191 # The weights and biases of layer 1 are W1 and b1; and in general the
192 # weights and biases of layer i are Wi and bi. The
193 # biases are initialized to zero and the weights are initialized
194 # so that each parameter has mean 0 and standard deviation weight_scale.
195 #
196 # BATCHNORM: Initialize the gammas of each layer to 1 and the beta
197 # parameters to zero. The gamma and beta parameters for layer 1 should
198 # be self.params['gamma1'] and self.params['beta1']. For layer 2, they
199 # should be gamma2 and beta2, etc. Only use batchnorm if
self.use_batchnorm
200 # is true and DO NOT do batch normalize the output scores.
201 # ===== #
202 dims = np.hstack((input_dim, hidden_dims, num_classes))
203
204 for i in range(1, self.num_layers+1):
205     Wi = 'W' + str(i)
206     bi = 'b' + str(i)
207     self.params[Wi] = weight_scale * \
208         np.random.randn(dims[i-1], dims[i])
209     self.params[bi] = np.zeros(dims[i])
210
211     if self.use_batchnorm:
212         if i != self.num_layers:
213             gammai = 'gamma' + str(i)
214             betai = 'beta' + str(i)
215             self.params[gammai] = np.ones(hidden_dims[i-1], )
216             self.params[betai] = np.zeros(hidden_dims[i-1], )
217
218 # ===== #
219 # END YOUR CODE HERE
220 # ===== #
221
222 # When using dropout we need to pass a dropout_param dictionary to each
223 # dropout layer so that the layer knows the dropout probability and the mode
224 # (train / test). You can pass the same dropout_param to each dropout layer.
225 self.dropout_param = {}
226 if self.use_dropout:
227     self.dropout_param = {'mode': 'train', 'p': dropout}
228     if seed is not None:
229         self.dropout_param['seed'] = seed
230
231 # With batch normalization we need to keep track of running means and
232 # variances, so we need to pass a special bn_param object to each batch
233 # normalization layer. You should pass self.bn_params[0] to the forward pass
234 # of the first batch normalization layer, self.bn_params[1] to the forward
235 # pass of the second batch normalization layer, etc.
236 self.bn_params = []
237 if self.use_batchnorm:
238     self.bn_params = [{'mode': 'train'}
239                       for i in np.arange(self.num_layers - 1)]

```

```

240
241     # Cast all parameters to the correct datatype
242     for k, v in self.params.items():
243         self.params[k] = v.astype(dtype)
244
245 def loss(self, X, y=None):
246     """
247     Compute loss and gradient for the fully-connected net.
248
249     Input / output: Same as TwoLayerNet above.
250     """
251     X = X.astype(self.dtype)
252     mode = 'test' if y is None else 'train'
253
254     # Set train/test mode for batchnorm params and dropout param since they
255     # behave differently during training and testing.
256     if self.dropout_param is not None:
257         self.dropout_param['mode'] = mode
258     if self.use_batchnorm:
259         for bn_param in self.bn_params:
260             bn_param[mode] = mode
261
262     scores = None
263
264     # ===== #
265     # YOUR CODE HERE:
266     # Implement the forward pass of the FC net and store the output
267     # scores as the variable "scores".
268     #
269     # BATCHNORM: If self.use_batchnorm is true, insert a bathnorm layer
270     # between the affine_forward and relu_forward layers. You may
271     # also write an affine_batchnorm_relu() function in layer_utils.py.
272     #
273     # DROPOUT: If dropout is non-zero, insert a dropout layer after
274     # every ReLU layer.
275     # ===== #
276     cache = {}
277     dropout = self.dropout_param.get('p')
278     use_dropout = dropout is not None
279     h = np.copy(X)
280     #print('self.num_layers: ', self.num_layers)
281     for i in range(1, self.num_layers):
282         Wi = 'W' + str(i)
283         bi = 'b' + str(i)
284         ci = 'c' + str(i)
285         if self.use_batchnorm:
286             gammai = 'gamma' + str(i)
287             betai = 'beta' + str(i)
288             h, fc_cache = affine_forward(
289                 h, self.params[Wi], self.params[bi])
290             h, bn_cache = batchnorm_forward(
291                 h, self.params[gammai], self.params[betai], self.bn_params[i-1])
292             h, relu_cache = relu_forward(h)
293             cache[ci] = (fc_cache, bn_cache, relu_cache)
294         else:
295             h, cache[ci] = affine_relu_forward(
296                 h, self.params[Wi], self.params[bi])
297
298     if use_dropout:
299         h, dropout_cache = dropout_forward(h, self.dropout_param)

```

```

300         cache[ci] = *cache[ci], dropout_cache
301     #print('self.params', self.params)
302
303     scores, cache['c' + str(self.num_layers)] = affine_forward(h,
self.params['W' + str(self.num_layers)],
304 self.params['b'+str(self.num_layers)])
305
306     # ===== #
307     # END YOUR CODE HERE
308     # ===== #
309
310     # If test mode return early
311     if mode == 'test':
312         return scores
313
314     loss, grads = 0.0, {}
315     # ===== #
316     # YOUR CODE HERE:
317     # Implement the backwards pass of the FC net and store the gradients
318     # in the grads dict, so that grads[k] is the gradient of self.params[k]
319     # Be sure your L2 regularization includes a 0.5 factor.
320     #
321     # BATCHNORM: Incorporate the backward pass of the batchnorm.
322     #
323     # DROPOUT: Incorporate the backward pass of dropout.
324     # ===== #
325     loss, ds = softmax_loss(scores, y)
326     dh = np.copy(ds)
327     dh, dw, db = affine_backward(dh, cache['c' + str(self.num_layers)])
328     grads['W'+str(self.num_layers)] = dw
329     grads['b'+str(self.num_layers)] = db
330
331     for i in range(self.num_layers-1, 0, -1):
332         Wi = 'W' + str(i)
333         bi = 'b' + str(i)
334         ci = 'c' + str(i)
335
336         if use_dropout:
337             dropout_cache = cache[ci][-1]
338             dh = dropout_backward(dh, dropout_cache)
339
340         if not self.use_batchnorm:
341             _Cache = cache[ci][0], cache[ci][1]
342             dh, dw, db = affine_relu_backward(dh, _Cache)
343         else:
344             fc_cache = cache[ci][0]
345             bn_cache = cache[ci][1]
346             relu_cache = cache[ci][2]
347             dh = relu_backward(dh, relu_cache)
348             dh, dgamma, dbeta = batchnorm_backward(dh, bn_cache)
349             grads['gamma'+str(i)] = dgamma
350             grads['beta'+str(i)] = dbeta
351             dh, dw, db = affine_backward(dh, fc_cache)
352
353             grads[Wi] = dw
354             grads[bi] = db
355             grads[Wi] += self.reg * self.params[Wi]
356             loss += np.sum(self.params[Wi]**2) * self.reg / 2
357     # ===== #

```

```
358 # END YOUR CODE HERE
359 # ===== #
360
361 return loss, grads
362
```